

# Limits of Metamorphic Testing for Object-Level Authorization: Temporal Reachability, Oracle Admissibility, and Measurement Resolution

Qiulong Lin<sup>a,b</sup>, Nanlei Chen<sup>a,b,\*</sup>

<sup>a</sup>Yango University, Fuzhou, 350015, China

<sup>b</sup>Sogang University, Seoul, 04107, Republic of Korea

## ARTICLE INFO

### Keywords:

oracle problem  
metamorphic testing  
object-level authorization  
BOLA/BFLA  
specificity measurement  
false proof

## ABSTRACT

Metamorphic testing is an established approach to the Web-security oracle problem: MST-wi defines 76 system-agnostic relations (MRs) over a Web-specific domain language, reporting 99.81% specificity. For the reachability-gated authorization relations examined here, the predicates delegate adjudication to a *GUI-reachability model* that is a *crawl-time snapshot*, and that delegation fails deterministically, in one direction.

We obtain three findings from a three-phase instantiation of the predicates as defined—crawl, out-of-band authorization change, execute—a 576-record scenario/configuration study, and a Gitea 1.22.6 study replicated on Gogs 0.14.3 and GitLab 17.11.7. (i) Access legitimately granted out of band after the crawl is reported as a *proven violation*—a *false proof*—because one of its two representations is unavailable. (ii) Repairing this with a readable, sound oracle that fully resolves the proposition leaves the MR step adding no adjudication information: a necessary condition, since a direct consumer of the same oracle and outcome reproduces the verdict, and an ordinary request changes nothing measurable (a descriptive point estimate, not a non-inferiority test, and not domination). (iii) The reported specificity is not conditioned on a reachability-authorization check, and its input-level denominator cannot resolve a scenario-level failure. The study instantiates 2 reachability-gated relations; the wider catalogue is support evidence.

On Gitea the mismatch appears only under the crawl-time reading, never live: it tracks the reading time, not visibility. The oracle that would discharge the condition is readable from an owner's vantage but never a black-box subject's, and executing the published relations reproduces the false proof within the gated subset that fires.

## 1. Introduction

Broken object-level authorization (BOLA), first in the OWASP API Security Top 10 [1], is the failure to enforce, per object, that the authenticated subject is entitled to the object touched. Detecting it is an *oracle* problem: vulnerable and correct systems answer with the same success status. Metamorphic testing is an established approach to that oracle problem, and MST-wi [2] is a substantial system-agnostic instance for Web systems: 76 system-agnostic relations (MRs) over 55 Web-specific predicates, reported to detect 85% of vulnerabilities with specificity 99.81%.

This paper asks what happens to that claim when the asserted property is object-level authorization. The predicates do not express entitlement; they express what their authors call reachability—what “the functionality exposed by the user interface” permits—with the role-to-resource mapping “**not always available**” [2]. We take that sentence as premise, not discovery. Delegating adjudication to a reachability model makes the moment the model was read part of the specification, so we instantiate the predicates *as defined*—over a model recorded at crawl time—and drive the system through three phases: crawl, change authorization out of band, execute. The middle phase is the one the original evaluation does not exercise.

We ran a 576-record scenario/configuration study on a constructed system and a controlled study on Gitea 1.22.6, replicated on Gogs 0.14.3 and GitLab 17.11.7. We propose no mechanism. What survives is three findings (Section 6), one per research question:

*Finding 1 (RQ1, temporal).* The adjudication as designed produces *deterministic false proofs*—a legitimate access granted out of band after the crawl is reported as a proven violation, at rate 1.00 on the constructed system and 0.4000 on Gitea, because the predicate's two representations cannot describe an out-of-band grant; the error is a property of the *reading*, not of the system.

*Finding 2 (RQ2, oracle).* When a sound, readable authorization oracle is available, the metamorphic step carries no adjudication information the oracle-plus-request pair does not—a necessary condition (Lemma 1), not an empirical surprise, whose empirical face is that replacing the step with a single ordinary request changes nothing measurable (false proofs 0 vs. 0, recall 0.50 vs. 0.49; a descriptive, point-estimate comparison, not a formal hypothesis test, no margin pre-specified, not domination), and on all 3 disagreeing records the MR-based method errs.

*Finding 3 (RQ3, measurement).* The reported specificity is not conditioned on an explicit check that the reachability model coincides with the authorization state, and is computed in units insensitive to scenario-uniform failure.

\*Corresponding author.

✉ ubfbzu3@163.com (N. Chen)

The native execution, the seeded target and the second implementation are support evidence for these findings, not further headline findings (Section 7).

*Research questions* The three findings answer one question each:

- **RQ1 (temporal).** Does the crawl-time reachability model remain sound for authorization decisions after an out-of-band authorization change?
- **RQ2 (oracle).** When a readable, sound authorization oracle fully resolves the proposition under test, does the metamorphic step add incremental adjudication information beyond the oracle and the request outcome?
- **RQ3 (measurement).** Can an input-level specificity metric provide stable scenario-level resolution for a scenario-uniform failure, or can the failure be arbitrarily diluted by unaffected inputs?

Section 6 answers each in turn.

The findings are not a claim that MST-wi is wrong: the reachability stand-in is announced by its authors, and the predecessor design attributes its residual false alarms to crawler coverage [3]. Our contribution is therefore a *methodological characterisation*, not a new detector, and has three layers: (i) *mechanistic characterisation*—we locate the false-proof failure structurally, in *when* the reachability model was read rather than which objects are visible (Finding 1); (ii) *fair-control ablation*—under a readable, sound authorization oracle, the metamorphic step adds no measurable adjudication information beyond the oracle and the request outcome, established by a same-observable direct adjudicator (Lemma 1, Finding 2); and (iii) *metric characterisation*—we formalise the input-level specificity denominator’s resolution limit for a scenario-uniform failure (Finding 3, Proposition 1). The evidence—the three-phase instantiation with internal controls (Section 3), the scenario/configuration study (Section 4), the Gitea study with an endogenous control and a seeded-vulnerability positive control (Section 5), the source reading, the native execution of the library’s own relations, and the second implementation (Section 7)—supports these layers rather than standing as further headline claims.

## 2. Background: the predicates in question

For authorization there is no intrinsic oracle: vulnerable and correct executions both answer with success. Metamorphic testing replaces the absent oracle with a relation over multiple executions—and the relation is only as sound as the assumptions it needs, which for authorization are exactly the facts the oracle was supposed to supply. MST-wi [2], building on Mai et al. [3], crawls a per-user model of reachable URLs and contents, then checks MRs against that recorded model. In the artifact we inspected, the predicates guarding an authorization MR read either the recorded model

directly (`cannotReachThroughGUI`, `userCanRetrieveContent`) or static subject properties; we call a predicate *GUI-derived* when its value is a function of the crawl.

*Why the false proof is a design consequence, not a bug* A single example makes the mechanism concrete. Suppose a subject can reach an object at crawl time, so the record stores that object as reachable. Afterwards an administrator grants the same subject access to a *further* object through a channel—an administrative endpoint or a share link—that a crawl of the subject’s own interface never traverses. At execution the guard still reads the frozen record, finds the new object absent from it, and concludes the subject cannot reach that object; the relation then treats the successful response as a violation. Every step is exactly as specified, so no implemented guard is at fault and no individual relation is miswritten. The false proof is structural: the predicate’s two representations—“reachable as recorded” and “not reachable as recorded”—cannot express “reachable, but only after the record was taken”. More crawling does not remove it, because the grant need not be visible through the subject’s interface at all; and the reported specificity never registers it, because the error is a property of the adjudication’s *reading moment* rather than of any single generated input. This is the mechanism the rest of the paper measures, and the reason we speak of a boundary of the specification rather than a defect of an implementation.

Two mechanical counts frame what follows. Over the 76 catalog MRs, 38 (50.0%) gate on a GUI-derived or identity predicate, and only 14 (18.42%) are authorization-relevant by name or action—an upper bound, reported as a caveat. And a full-text search for predicates reading the *state of an authorization artifact* (ownership, permission record, share) finds 0 occurrences. Within the examined predicate vocabulary, the authorization MRs therefore have no explicit authorization-artifact state observable on which to base the decision, and the reading moment enters the specification: if an entitlement changes after the crawl, the MR evaluates reachability against a world that no longer exists. This is a statement about what the relations *mean*; our experiments are built so that the crawl-time and live readings can be told apart.

## 3. Protocol: a faithful three-phase instantiation

Two instruments carry distinct claims. The three-phase instantiation below carries the quantitative results of Sections 4 and 5; a direct execution of the library’s own relations—source unmodified, with the build configuration and a replay harness added—carries one qualitative claim, that the published code produces the false proof we predict (Section 7.2). Neither instrument is a benchmark: we report no sensitivity or specificity attributable to MST-wi.

*Three phases* (1) *Crawl*: record, per subject and object, whether the subject can reach the object through the interface, and freeze that record—the model the predicates

will consult. (2) *Change authorization out of band*: alter a (subject, object) entitlement through a channel outside the crawled interface; the frozen model is not updated, because the architecture has no mechanism to update it. (3) *Execute the MR* against the current state while the predicates consult the frozen model. Phase 2 is what the original evaluation does not exercise, and it is where the reachability stand-in can be wrong.

*Adjudicator variants and internal controls* A raw alarm becomes a verdict through an exception policy, and different policies are different methods: **v1** (raw alarms proven); **v2** (legitimate-sharing exceptions filtered); **v3** (state-calibrated—consults a readable authorization introspection surface, abstains when none is readable); `degrade_all` (the floor: abstain on every alarm); `direct_only` (authorization surface only, access outcome ignored—a deliberately unfair control); and `direct_observe` (a succeeding plain request checked against the same readable surface). The last consumes the same observable quantity as the MR family and decides the negative result.

*Ground truth* Every quantitative result names its ground-truth source, and it is never the detector under test: construction state for the constructed system and the scenario/configuration study; re-probed live authorization (the grantee’s response hash matched against the owner’s, with a third account refused) for Gitea and the cross-vendor systems; the pinned route tables for the grant census; and the seed list for the seeded target.

*Scenarios and design cells* Seven scenarios separate the reasons an MR can misfire: no shared object (S1, S2); an in-band legitimate share visible in the framework log (S3); an unshared plain violation (S4); a read-only share plus a write violation (S5); an out-of-band share invisible to the framework log (S6); an out-of-band share plus a genuine violation (S7). They are crossed with 8 configurations defined by instantiation mode, enforcement correctness, channel visibility and introspection availability. The nominal cross is  $7 \times 8 = 56$  combinations, but four are inadmissible: the two no-exception configurations (`V-noE3-vuIn0`, `V-noE3-vuIn1`) are undefined under the two out-of-band scenarios (S6, S7), where the legitimate-sharing exception does not apply. This leaves 52 admissible design cells.

*Statistical treatment* The 576 records are not independent observations. Each admissible cell is expanded deterministically by the factors that are actually varied: 2 metamorphic relations (MR-002, MR-004), 3 exception-policy variants (raw, one level, two levels—the first deliberately thrown switch), and 2 reachability readings: the frozen crawl-time snapshot and the live execution-time state—the second deliberately thrown switch). The two relations do not both apply to every cell: only MR-002 applies under scenario S5, so 8 S5 cells carry one relation while the remaining 44 carry

two. The record count is therefore

$$576 = 44 \cdot (2 \cdot 3 \cdot 2) + 8 \cdot (1 \cdot 3 \cdot 2),$$

a deliberately unbalanced design—the per-cell multiplier differs by scenario. The records are therefore a fixed enumeration, not a sample: we report counts, rates and design-cell denominators, and we attach no inferential interval to any of them. Where a headline count has a small denominator we say so explicitly (the anchored rate,  $2/2$ , is unstable by construction) and lead with the larger-denominator count of the same phenomenon. No conclusion rests on statistical significance. Where a comparison is paired—the ablation of Section 4.3—we report discordant-pair counts rather than an interval over a difference. The replication ceiling is explicit: one constructed system, 52 cells, 2 MRs, three real systems.

*Tooling and code provenance* Development of the analysis scripts was assisted by an AI coding assistant (declared in Section 10). Every script that assistant touched was reviewed, executed, and its output checked by the authors against the raw records; the released pipeline recomputes every reported scalar from those records with no hand-entered value in the body. No AI tool participated in the detection loop or in any adjudication.

## 4. Synthetic study: where the adjudication goes

This study supplies the evidence for Findings 1 and 3, and the ablation behind Finding 2; the findings themselves are stated in Section 6.

### 4.1. Definitions

Confusions are defined over the *verdict*: a *false proof* is the strict subset of false alarms promoted to violation-proven—an error the tool asserts as established. We also report the *detection-keep rate* (share of fired true violations still alarmed) and count only abstentions justified by an unreadable oracle.

### 4.2. What the constructed system shows

*Attribution is clean* On the anchored configuration the false-proof rate is  $2/2 = 100.0\%$ , but the denominator is 2, so the rate is unstable by construction and we do not attach an interval to it. With the legitimate-sharing exception absent the rate is 0.00 (0/2), and with the channel listed it is 0.00 (0/2). The phenomenon is therefore attributable to a legitimate sharing exception the MR cannot know about, not to execution itself. A  $2 \times 2$  grid over channel visibility and enforcement correctness tells the same story from the other side: on unlisted cells the false-proof rate is 44.44% with no false negatives; on listed cells it drops to 0.0% but a false negative appears among the five violating records (1 of 4+1). Avoiding false proofs and keeping detection are controlled by one switch, and the switch cannot be thrown both ways at once.

**Table 1**

End-to-end confusion matrix over all 576 records (216 true violations, 360 correct cases). The records are a fixed enumeration, not a sample (Section 3), so no inferential interval is attached; where a denominator is small we say so rather than interval it. `direct_only` is a deliberately unfair control, shown to make that point.

| method                | TP  | FP  | FN  | TN  | precision | recall | F1   |
|-----------------------|-----|-----|-----|-----|-----------|--------|------|
| v1 (raw)              | 198 | 60  | 18  | 300 | 0.77      | 0.92   | 0.84 |
| v2 (exception-aware)  | 213 | 60  | 3   | 300 | 0.78      | 0.99   | 0.87 |
| v3 (state-calibrated) | 105 | 0   | 111 | 360 | 1.00      | 0.49   | 0.65 |
| degrade_all           | 0   | 0   | 216 | 360 | —         | 0.00   | —    |
| direct_only           | 108 | 108 | 108 | 252 | 0.50      | 0.50   | 0.50 |
| direct_observe        | 108 | 0   | 108 | 360 | 1.00      | 0.50   | 0.67 |

*Filtering exceptions does not help* The exception ladder—raw, one exception level, two—leaves the anchored rate at 1.00, 1.00, 1.00: it does not move. The exception levels we can express are themselves functions of the same reachability model, so they inherit its blind spot.

*The framework cannot see the channel that matters* When a legitimate share was created through the framework’s facility, its log holds 12 records and v2 produces 0 false proofs (rate 0.00); when the same share was created out of band, the log holds 0 records—no evidence the share exists—and v2’s false proofs rise to 8 of 8 alarms (rate 1.00), every one promoted to violation-proven. The MR cannot distinguish “this share does not exist” from “this share exists but was not crawled”; its only representation of the first is the second. On S7, where a genuine violation *does* sit behind the out-of-band share, the same blindness makes the alarm right for the wrong reason (6 of 6 reached)—and that coincidence of alarms with violations is what an input-level metric cannot unpack. The same tension appears within a single MR: for S3 (must not alarm) vs S5 (must alarm), no configuration in the family is correct on both—the choice of switch is a choice of *which* error to make.

*Stratification, not aggregation* Aggregated, v3 looks best (recovery 1.00 against degrade\_all’s keep rate 0.00); the reading is wrong, because v3 changes regime with oracle readability. Stratified: on the readable stratum (36 records) it removes v2’s false proofs (0 vs. 4) while still alarming on all 9 fired true violations; on the absent stratum (18) it alarms on none of the 5 true violations v2 still alarms on, abstaining 13 times. Pooling converts honest abstention into apparent performance, so v3’s numbers may appear only stratified.

#### 4.3. The decisive ablation: does the metamorphic step carry work?

Our earlier iteration proposed v3 as a mechanism, resting on the claim that the metamorphic step was necessary. We tested that claim by removing the step and substituting the most naive consumer of the same observable. The distinction matters: `direct_only`—oracle yes, access outcome ignored—produces 108 false proofs against v3’s 0 (precision 0.50), and rejecting it proves nothing because it discards half the evidence. The fair control is `direct_observe`: a succeeding

ordinary request plus the same readable surface, consuming exactly what the MR-based adjudicator consumes.

Table 1 closes this line of work. `direct_observe` has 0 false proofs against v3’s 0—a tie at zero, not a win. The comparison is paired, so we report the discordance rather than an inferential interval: over the 216 true-violation records the two methods disagree on 3, and on every one of them `direct_observe` is right—it detects a violation that v3 misses or abstains on, its reachability precondition blocking a detection the plain request makes. The aggregate recalls 0.50 and 0.49 restate that asymmetry rather than measure a difference, and no margin was pre-specified, so no formal hypothesis-test claim attaches. We found no incremental adjudication advantage of the metamorphic step over direct observation in these oracle-available scenarios, and that absence provides no support for an incremental adjudication benefit in these scenarios: the step would have to be better at something the plain request cannot do. Selective presentation would have hidden the cost: v3’s precision of 1.00 was bought with 111 false negatives (F1 0.65), an exchange of 60 false alarms for missed violations dressed as an improvement.

#### 4.4. Input-level specificity and scenario-level resolution

Under the input-level specificity convention used by MST-wi [2], the reported specificity has the form

$$S_{\text{input}} = 1 - \frac{\text{false alarms}}{\text{inputs}}.$$

We cannot reproduce the engine’s exact input count, so we approximate two follow-up inputs per case—enough to exhibit the denominator’s resolution, the only point at issue. In that approximation the configuration in which a whole class of legitimate access is misreported still yields a specificity of 0.846154 (84.62%). A failure uniform across a scenario class is diluted by every other input in the run, and the predecessor design says the same of itself, attributing its residual false alarms to crawler coverage of “all URLs for all users” [3].

*A proposition on the metric’s resolution* Define  $S_{\text{input}} = 1 - FP_{\text{inputs}}/N_{\text{inputs}}$  and  $S_{\text{scenario}} = 1 - N_{\text{failed}}/N_{\text{scenarios}}$ , where  $FP_{\text{inputs}}$  counts false alarms over all generated inputs and  $N_{\text{failed}}$  the scenarios in which the failure is uniform.  $S_{\text{scenario}}$  is introduced only as an analytical construct for

characterising resolution; it is not proposed as a replacement evaluation metric.

**Proposition 1.** Fix  $f \geq 1$  failed scenarios, each producing a fixed number  $a \geq 1$  of false-alarm inputs, and  $s \geq 1$  successful scenarios each contributing  $m$  additional, alarm-free inputs. Then  $S_{\text{input}} = 1 - (fa)/(fa + sm)$ , which tends to 1 as  $m \rightarrow \infty$ , while  $S_{\text{scenario}} = 1 - f/(f + s)$  is independent of  $m$  and strictly less than 1.

*Proof.*  $N_{\text{inputs}} = fa + sm$  and  $FP_{\text{inputs}} = fa$ ; substitution gives  $S_{\text{input}} = 1 - (fa)/(fa + sm) \rightarrow 1$  as  $m \rightarrow \infty$ , whereas the scenario-level figure depends only on the fixed counts  $f$  and  $s$ .  $\square$

**Corollary 1.** No fixed input-level specificity threshold  $t < 1$  can guarantee that a scenario-uniform failure will push the metric below it: choosing  $m > (fa/s) \cdot t/(1 - t)$  drives  $S_{\text{input}}$  above  $t$  while the  $f$  failing scenarios remain failing.

The metric is therefore not merely imprecise: it is arbitrarily dilutable with respect to a scenario-uniform failure—adding unaffected inputs drives it toward 1 without touching the failure—so it lacks stable scenario-level resolution.

*Implementation robustness* An independently written implementation of the same predicates was compared on a subset fixed in advance; it bounds implementation accident only, and is reported with the confirmations in Section 7.3.

## 5. Real systems: Gitea, Gogs, GitLab, and a seeded control

This section supplies the real-system evidence for Findings 1 and 2: the Gitea study with its endogenous control, the cross-vendor replication, and the seeded positive control (support evidence, Section 7).

A synthetic system cannot tell us whether the crawl-time reading we identified is an artefact of our construction, so we repeated the essential comparison on real systems. We chose Gitea 1.22.6 for genuine object-level sharing with per-object visibility, a machine-readable sharing surface, and a documented administrative grant path. Three scope conditions on the experiment's conclusions, stated before the numbers. First, on the endpoints and authorization transitions exercised in this study, Gitea enforced the expected decisions: no instance of the studied vulnerability was found, so the experiment measures only the false-proof side. Second, the cell study uses our instantiation of the predicates, not the engine; the engine's own relations are executed separately in Section 7.2. Third, the reachability model is reported under *both* readings—crawl-time snapshot (what the definitions presuppose) and the live execution-time state—so the explanatory variable, the time at which the model was read, can be isolated. The experiment produces 96 cells, 48 per reading.

**Table 2**

The grant surface of Gitea 1.22.6 (instance-independent).

| quantity                              | count |
|---------------------------------------|-------|
| paths in the system's API description | 253   |
| grant endpoints                       | 15    |
| revocation endpoints                  | 2     |
| of which usable without site-admin    | 8     |

**Table 3**

Authorization-introspection readability protocol on Gitea. Each row is a distinct population and denominator, so the rows are not comparable: both the vantage and the unit of the denominator differ.

| vantage                                        | attempts | readable |
|------------------------------------------------|----------|----------|
| owner-vantage                                  | 56       | 56       |
| authorization-introspection surface            |          |          |
| arbitrary subject (black box)                  | 40       | 0        |
| subject also owner/admin (representative grid) | 56       | 16       |

*Grant surface* Table 2 counts an endpoint as grant-capable if, without modifying the system, a request can move a (subject, object) pair from not-entitled to entitled. The 8 endpoints drivable by a non-administrator are the out-of-band channel in its real form: a subject who is not an administrator can create an entitlement that a crawl of that subject's own interface need not reveal. The cross-vendor table below reports a normalised census (semantically equivalent routes merged, non-authorization grant surface excluded), which is why its Gitea figure (13) differs from the raw count (15) here; the two answer different questions.

*Is the authorization surface readable, and from where?* Readability depends entirely on the vantage, so the counts below use different populations and must not be read as one denominator. From the owner/administrator vantage the sharing surface is readable on 56 of 56 attempts: complete. From a black-box subject's own vantage—the vantage a black-box tool occupies—it is readable on 0 of 40: never. An intermediate count of 16/56 comes only from subjects who happen to be owners or administrators, and does not constrain the mechanism; it is therefore not synonymous with “an arbitrary black-box subject can read the surface”. Table 3 collects these populations. On 48 of 48 applicable cells the credential-substitution check confirmed that our model agrees with the system's behaviour.

*The mismatch is in time, not in visibility* Table 4 is the load-bearing result: the adjudication departs from the real authorization state on 10 of 48 cells under the crawl-time reading and 0 of the same cells under the live reading. The system state and the tested authorization relation were held fixed; only the time at which the reachability model was read

**Table 4**

Departure from the real authorization state, per reading of the reachability model.

| reading of the reachability model                   | mismatches | cells |
|-----------------------------------------------------|------------|-------|
| crawl-time snapshot (as the definitions presuppose) | 10         | 48    |
| live execution-time state                           | 0          | 48    |

**Table 5**

Endogenous control: two objects that are both public at the end.

| object                      | cells | cannotReach true | MR fires |
|-----------------------------|-------|------------------|----------|
| public before the crawl     | 6     | 0                | 0        |
| public only after the crawl | 6     | 5                | 5        |

was changed. Under the crawl-time reading, a bookkeeping adjudication produces 10 false proofs on the 25 authorized cells (rate 0.4000), against 0 alarms under the live reading; v3 produces 0 false proofs, and `direct_observe` emits 0. The unfair control `direct_only` raises 23 alarms, more than any other adjudicator, simply because it ignores the request outcome; on this system none of them is promoted to violation-proven under our strict false-proof definition (0 false proofs), so the alarm count and the false-proof count must not be read as the same quantity. Channel provenance of the authorized access (24 of 25 not observable by the framework log, 0.9600) is a property of our topology—a design parameter, not a field prevalence.

*An endogenous control* The most serious alternative explanation is that public objects are simply harder to enumerate—visibility, not time. Two objects that end equally public but became public at different moments control for this (Table 5): the predicate reports unreachability on 5 of 6 cells only for the object made public after the crawl, and on 0 for the object public throughout, with the MR firing correspondingly (5 vs. 0). Final visibility is identical; within this controlled design, the only manipulated explanatory variable is when the model was read.

### 5.1. Cross-vendor replication: Gogs 0.14.3 and GitLab 17.11.7

We repeated the same three-phase protocol and 48-cell design on two further systems (Table 6): **Gogs** 0.14.3, the project Gitea forked from in 2016 (related lineage, independently evolved code; census from its pinned route table, 80 paths), and **GitLab** 17.11.7 (no lineage; census from the pinned running instance’s routes, 1357 paths, since its OpenAPI document omits core grant endpoints). Gogs is a replication across implementations, GitLab across vendors within one domain.

The replication changes one thing and confirms two. What it changes concerns the surface: *readability does not imply soundness*. Gitea’s owner-vantage surface is readable and sound (agrees with effective access on all 48 cells). Gogs’s collaborator-membership endpoint is

equally readable (48/48) but unsound—wrong on 17 cells (35.4%), exactly the authorizations arising outside direct membership: site administration (8), public visibility (5), out-of-band grants (4). GitLab’s `members/all/{u}` resolves inherited membership but public visibility and site administration confer no membership: 18 false negatives (37.5%). What the surface cannot see is predictable from what it measures—why clause 2 of the admissibility condition (Section 6.4) cannot be automated from readability alone. The consequence is measurable: on Gitea the surface-gated adjudicators produced no false proofs (0 v3, 0 direct-observe), while on Gogs they manufacture them (8 snapshot, 5 live, 17 direct-observe) and on GitLab again (5, 18), every v3 false proof falling on a visibility-channel cell. A readable membership surface with a blind spot converts that blind spot into confirmed violations; no admissible oracle was available from the tested introspection surfaces on Gogs or GitLab, so the honest adjudicator abstains on the stratum where those surfaces would license a promotion.

Two results replicate everywhere: bookkeeping alarms false-proof on legitimately granted access (10/11/10 cells) with the same out-of-band-dominant grant profile; and the reachability mismatch persists (10/19/10 snapshot). On Gogs the mismatch only partially resolves live (13 remain—team or public grants never enter a subject’s personal crawl); on GitLab the live reading recovers everything (0 remain)—a second, vendor-specific failure of the snapshot premise, in the crawl model rather than the oracle.

### 5.2. Support evidence: a positive control with a seeded vulnerability

On the tested endpoints and transitions the three real systems enforce the expected decisions, so the question remains: *when a real object-level flaw is present, does the same detector find it?* We answer with a positive control. A *controlled authorization target* faithfully emulates the Gitea API contract the engine consumes, with a correct internal authorization model and 3 seeded vulnerable instances—all instantiating one fault pattern (a none-entitlement subject served the object at `GET /api/v1/repos/{owner}/{repo}`),

**Table 6**

Cross-vendor replication under the identical three-phase protocol (crawl, out-of-band authorization change, adjudicate) on Gitea 1.22.6, Gogs 0.14.3, and GitLab 17.11.7. Counts are per reading of the same 48-cell grid; the live reading is given in parentheses where both exist. The snapshot reading is the operating point of the published relations. Gitea and Gogs share lineage (fork, 2016); GitLab shares none. The team-derived row is read out of the per-system oracle false negatives by channel (Section 5.1). The Gitea grant count here is the *normalised* M1 census (semantically equivalent routes merged, non-authorization grant surface excluded), and therefore differs from the raw route count in Table 2; the two tables use different normalisations and are not in conflict.

|                                                        | Gitea 1.22.6 | Gogs 0.14.3 | GitLab 17.11.7 |
|--------------------------------------------------------|--------------|-------------|----------------|
| API entries in pinned version (M1 denominator)         | 253          | 80          | 1357           |
| Grant endpoints (M1)                                   | 13           | 10          | 18             |
| of which usable without site admin                     | 10           | 4           | 17             |
| Owner-vantage oracle readable                          | 48           | 48          | 48             |
| oracle agrees with actual access                       | 48           | 31          | 30             |
| oracle wrong (false negatives)                         | 0            | 17          | 18             |
| oracle wrong: team-derived access missed               | 0            | 4           | 0              |
| Subject-vantage oracle readable                        | 8            | 35          | 48             |
| Reachability mismatch, snapshot (live)                 | 10 (0)       | 19 (13)     | 10 (0)         |
| False proofs: bookkeeping alarm, snapshot (live)       | 10 (0)       | 11 (5)      | 10 (0)         |
| False proofs: v3 (oracle-gated), snapshot (live)       | 0 (0)        | 8 (5)       | 5 (0)          |
| False proofs: direct-only (oracle only)                | 0 (0)        | 17 (17)     | 18 (18)        |
| False proofs: direct-observe (oracle $\wedge$ request) | 0 (0)        | 17 (17)     | 18 (18)        |

**Table 7**

Positive control: detection of 3 seeded vulnerable instances (one fault pattern). TP/FN/FP/TN are counted against the known seed, not against a model.

| method                   | TP | FN | FP | TN |
|--------------------------|----|----|----|----|
| mr_bookkeeping           | 3  | 0  | 2  | 43 |
| v3 (owner vantage)       | 3  | 0  | 0  | 45 |
| v3 (subject vantage)     | 3  | 0  | 0  | 45 |
| direct_only              | 3  | 0  | 30 | 15 |
| direct_observe           | 3  | 0  | 0  | 45 |
| direct_observe (subject) | 3  | 0  | 0  | 45 |

differing in subject and object, not in fault location; Table 7 is therefore evidence about this fault pattern, not three independent defect discoveries. Two design points: the target deliberately exposes its authorization state to the subject-side harness (mirroring the admissibility condition), which production Gitea does not—on the seeded target the subject-side harness can read the authorization state of every cell by design, whereas on production Gitea the oracle is readable from the owner/admin vantage (56/56) but from an arbitrary subject's vantage it is not (0/40); and the introspection surface reports the *correct* none permission, so the oracle stays sound and the defect is visible only as a discrepancy between oracle and request. We drove it with the same topology, mutation, crawl, introspection and adjudication code as above, differing only in the system under test; ground truth comes from the seed list, not from the detector.

The observation-gated adjudicator is consistent with the negative result in this controlled fault pattern (Table 7): `direct_observe` and `v3` catch all 3 planted instances with zero false positives; `direct_only`, which ignores the request

outcome, raises 30 alarms on cells where the surface says none and no access succeeds—alarms our strict definition does not promote. The control narrows the claim to this fault pattern: when a surface that fully resolves the authorization question is available, the metamorphic step is not necessary for detection here, because both `direct_observe` and `v3` catch all 3 instances. It does not generalise beyond one fault pattern, and the seeded BFLA defect (below) remains outside the read-oriented logic's coverage. The control also delimits the detector: we seeded one BFLA defect (non-privileged bob succeeding at `PUT . . . /collaborators/{u}`, HTTP 204), which the read-oriented logic does not cover—a known boundary, reported as such rather than as a detection failure. This is a controlled research artefact, not production Gitea; it is a sanity check on the detector's logic, not a validation of detection effectiveness on a population, and not a statement about Gitea's CVE posture. Its source, seed list and run script are in the reproducibility package at <https://github.com/ubfbuz3/JISA> (tag artifact-v1).

## 6. Three findings

The contribution is three findings, each a boundary on what the published relations can establish about object-level authorization, and each carried by the studies of Sections 4–5: a *temporal* boundary, an *oracle* boundary, and a *measurement* boundary. The confirmations the findings rest on—the predicates read at source, the published relations executed on a real system, and an independent implementation of the same predicates—introduce no further claim and are collected as support evidence in Section 7.

### 6.1. Finding 1 (RQ1): deterministic false proofs—the temporal boundary

*Taken as published—raw alarms, or alarms filtered by bookkeeping exceptions—the adjudication reports a legitimate access granted out of band after the crawl as an already-established violation, in the direction forced by the two-valued predicate encoding.* On the constructed system, rate 1.00 (8 of 8 alarms) against 0 for the same share created in band; on Gitea, 10 false proofs on 25 authorized cells (0.4000) under the crawl-time reading, against 0 alarms under the live reading—objects, entitlements and predicates held fixed, only the reading time changed—so the error is a property of the *reading*, not of the system (Table 4). The endogenous control (Table 5) holds final visibility fixed and isolates the moment at which the model was read. The direction is forced: the predicate has exactly two representations, and a grant created outside the crawled interface has no representation but “cannot reach”. The predecessor attributes its residual false alarms to crawler coverage [3]; we add that the failure is a function of *when* the model was read, which makes the natural oracle-based repair available, but shows no incremental adjudication benefit from the metamorphic step (Finding 2).

### 6.2. Finding 2 (RQ2): with a sound, readable oracle, the metamorphic step adds no adjudication information

*If the false proofs of Finding 1 are removed by consulting a readable introspection surface, it is the oracle that carries the work: for the authorization relations and scenarios evaluated here, once the oracle fully resolves the authorization proposition under test, replacing the entire metamorphic step with a single ordinary request—one that consumes the same readable oracle and access outcome—provides no additional measurable adjudication benefit.*

We state the reason rather than present it as a discovered effect, because it is a necessary condition on the claim, not an empirical regularity.

**Lemma 1** (oracle-sufficient adjudication). Let  $O$  be an authorization oracle that is *readable* (its answer can be obtained) and *sound* (its answer equals the true permission), and let  $R$  be the observed outcome of an ordinary request. Any adjudicator whose verdict is a function  $V = f(O, R)$  of  $O$  and  $R$  alone is extensionally equal to the direct consumer  $D := f(O, R)$ : the two return the same verdict on every input. Consequently a metamorphic step can be *shown* to be necessary only by exhibiting an observable it consumes that the direct consumer does not.

*Proof.* An adjudicator of the stated form and  $D$  are the same function applied to the same pair of arguments, so they agree pointwise; the last sentence is the contrapositive.  $\square$

The lemma concerns *adjudication* information only. It does not imply that a metamorphic relation cannot improve input generation, search-space reduction, or execution efficiency: a relation may guide which requests are issued, or

which gated subset is exercised, while carrying no verdict information beyond  $(O, R)$ . It is a requirement on *claims*, not a verdict on metamorphic testing: it says only that a verdict computable from  $(O, R)$  cannot be shown to beat the direct consumer of  $(O, R)$ . The evaluated relations do consult a third quantity—the crawl-time reachability model  $M$ —so the empirical question is whether  $M$  carries authorization information that a readable, sound  $O$  does not already supply. For the authorization proposition under test, once  $O$  is readable, sound, and complete,  $M$  cannot supply additional decision-relevant information: agreement with the true state is redundant, while disagreement makes the resulting authorization verdict unsound (Finding 1). Source inspection confirms the two relations read only  $M$  and the response (Section 7.1); the experiment confirms the consequence.

Evidence: Table 1—0 false proofs vs. 0 (a tie at zero, not a win) over the 216 true-violation records. The comparison is paired, so we report the discordance rather than an inferential interval: the two methods disagree on 3 records, and on every one of them *direct\_observe* is right—the aggregate recalls 0.50 and 0.49 restate that asymmetry rather than measure a difference, and no margin was fixed in advance. Gitea independently reproduces the zero-false-proof result (0 vs. 0), but does not provide the same comparison because the tested surface contains no seeded vulnerability; the cross-vendor systems (§5.1) and the seeded control (§5.2) complete the evidence. One qualification is part of the claim: the control is *constructed*, not discovered—*direct\_observe* instantiates exactly the  $D$  of the lemma, so the tie is the empirical face of the necessary condition, a statement about the definitions, not a coincidental measurement.

### 6.3. Finding 3 (RQ3): the reported specificity lacks scenario-level resolution

*The reported 99.81% is a statement about a corpus of generated inputs; it is not conditioned on an explicit check that the reachability model coincides with the true authorization state, and its input-level denominator cannot provide stable scenario-level resolution for a scenario-uniform failure.* The omission: adjudication is a function of the reachability model, so where model and true state diverge—the out-of-band case of Finding 1—the figure measures agreement with the tool’s own crawler; the authors concede the mapping is “not always available” [2], but the number is reported without the condition. The denominator: a failure uniform across a scenario class contributes only as many false positives as that class emits inputs, and is diluted by the rest of the run (Proposition 1: in our approximation the configuration that misreports a whole class of legitimate access still scores 84.62%; Corollary 1: no fixed threshold can guarantee that such a class forces the metric below it). We are not saying the figure is false or the false positives misdiagnosed; the claim is that such a metric does not itself validate the reachability-to-authorization correspondence on which its interpretation depends, and that a class-uniform failure can be held above any fixed threshold by unaffected inputs.

#### 6.4. Admissibility, and falsification

Three properties must be kept apart. A surface is *readable* when its answer can be obtained from the vantage in question; *sound* when its answer equals the true authorization state; and *admissible*, for promotion to violation-proven, when it is readable *and* one can give a reason that it is not of the same kind as the defect under test—not itself a reachability artefact whose blind spot correlates with the blindness tested. Readability is what makes soundness usable, but it neither implies soundness nor confers admissibility: the owner-vantage authorization-introspection surface is readable on every cell in each system (48 Gitea, 48 Gogs, 48 GitLab), yet sound on Gitea alone—it agrees with effective access on 48, 31 and 30 of those cells respectively. The cross-vendor systems are therefore not cases of “no oracle”: a readable membership surface exists on both, but it is unsound for the authorization channels exercised here, so it fails clause (2). Clause 2 is a per-system judgement and not automatable from readability alone; the cross-vendor results make it concrete (Gogs and GitLab fail it; their adjudicators convert the blind spots into 8 and 5 false proofs). Where the condition holds, the honest method consumes the oracle directly and drops the relation; where it does not, the honest method abstains—as v3 does on the unreadable stratum, where the aggregate table would have ranked it best, a ranking that reflects abstention rather than improved detection. Each finding would fall as follows: Finding 1, if the predicates were defined over live rather than recorded reachability; Finding 2, by an authorization MR needing an observable that a succeeding request plus a readable oracle cannot produce; Finding 3, by a specificity in scenario units, or with the condition explicit, that still fails to separate the configurations studied here.

### 7. Support evidence

Nothing here is a further finding. This section holds what the three findings rest on: the semantics of the two governing predicates read at source (Section 7.1), an execution of the published relations on a real system (Section 7.2), and the robustness of our numbers to a second, independent implementation of the same predicates (Section 7.3); the seeded positive control (§5.2) and the cross-vendor replication (§5.1) are the remaining confirmations and stay with the study that produced them. Each item is a confirmation of a mechanism already stated, not a new claim about what the technique can detect.

#### 7.1. Confirmation at source

Definitions can be read generously, so we checked the predicates at source in the public library [2]. It ships 23 relations, 11 authorization relations, 8 gating on the two predicates analysed. Four implementation-level properties not made explicit in the main presentation: (i) *the freeze is hard-coded*—cannotReachThroughGUI is set-membership over the crawl inputs, issues no request, and is cached in a field written once and never invalidated (userCanRetrieveContent likewise reads a crawl-produced store loaded once); (ii) *the*

**Table 8**

Execution facts for the native run of the published relations (Section 7.2). All values are emitted by ingest\_native\_engine.py from the run record; upstream Java sources are unchanged apart from the build configuration.

| item                                                  | value       |
|-------------------------------------------------------|-------------|
| authorization relations: catalogue / executed / fired | 11 / 10 / 2 |
| engine alarms (the engine's own return value)         | 6           |
| HTTP requests attributed to observing relations       | 270         |
| changed upstream Java sources                         | 0           |
| build/config changes (one pom.xml line)               | 1           |
| added harness classes / added lines                   | two / 612   |

*clause that could have caught it is inert*—the “absent across all users” set is built by stopping each recorded sequence at the first login action, so it is always empty; (iii) *missing data defaults to an alarm*—with an empty or null crawl argument the predicate returns “cannot reach”, so an empty crawl *satisfies* the precondition: the mechanism behind Finding 1’s deterministic direction; (iv) *the role relation is configuration, not state*—the supervisor guard reads a static configuration file and reports true for identical accounts, so it cannot reflect a later grant.

#### 7.2. Executing the published relations

The library ships relations but no execution environment; we supplied the system configuration and crawl snapshot and ran the relations—source unmodified, build configuration (pom.xml) and an external harness added (1 changed upstream file, 0 changed Java sources, two added classes totalling 612 lines)—against Gitea 1.22.6 (Table 8). 10 of 11 authorization relations executed; 2 returned failures, 6 alarms in all—the engine’s own getFailures() value (Table 9). What the run establishes is scoped to the *reachability-gated subset that fires here*, not a full replay of the catalogue: one relation aborts on an engine-side defect, most issue no request because their own guards are unsatisfied, and the content predicate is never invoked. Almost none of the 1483.7 seconds is analysis: each leaf evaluation invokes a browser-cleanup hook costing minutes, a cost of the missing environment that bounds any reproduction.

The load-bearing cell: the subject bob holds a legitimate entitlement to alice/sl, granted out of band after the crawl. There the published cannotReachThroughGUI answers “cannot reach”—the answer it is *defined* to give, since the subject’s crawl record has no entry for the repository; across the run it answered so 1149 of 7321 invocations (15.7%), 5 on this cell, and 2 relations (OTG\_AUTHZ\_002, OTG\_AUTHZ\_002c) report a failure there. The access is authorized: all 3 granted cells were re-probed, and the grantee’s response hashes identically to the owner’s while a third account is refused. The 6 alarms fall on 3 cells, all post-crawl grants, none

**Table 9**

The published authorization relations executed against Gitea 1.22.6. Each row is the applicability chain from the catalogue to the outcome: *status* is the engine's per-relation result (10 of 11 evaluated, 1 aborted on an engine-side defect), *alarms* its own `getFailures()` return (2 relations, 6 in all), and *requests* (270 total) attributed to the relation observing them first.

| MR              | status   | alarms | requests |
|-----------------|----------|--------|----------|
| OTG_AUTHZ_001   | executed | 0      | 0        |
| OTG_AUTHZ_001b  | executed | 0      | 0        |
| OTG_AUTHZ_001b2 | executed | 0      | 0        |
| OTG_AUTHZ_002   | executed | 3      | 270      |
| OTG_AUTHZ_002a  | executed | 0      | 0        |
| OTG_AUTHZ_002b  | executed | 0      | 0        |
| OTG_AUTHZ_002c  | executed | 3      | 0        |
| OTG_AUTHZ_002d  | executed | 0      | 0        |
| OTG_AUTHZ_002e  | error    | –      | –        |
| OTG_AUTHZ_003   | executed | 0      | 0        |
| OTG_AUTHZ_004   | executed | 0      | 0        |

on a non-entitled object. Limits: the harness's HTTP observation channel is exact for the reachability predicate (set-membership, no request) but unexercised for the content predicate (0 invocations; supervisor guard 8642, 1321 true); not all gated relations fire, for reasons inside their own guards, and 1 abort on an engine-side defect—so the reproduced false proof is that of the reachability-gated subset that fires, not a full-replay count of the published relations. And Gitea, on the tested surface, enforces the expected decisions, so the run can only exhibit false proofs.

### 7.3. Implementation robustness

An independently written implementation of the same predicates was compared with the main one on a subset fixed before the comparison, not sampled after it: the in-band scenarios on which a second reading is defined, holding the exception-policy variant at the representative level and the reachability reading at the one each relation uses (MR-002 blind, MR-004 aware). This yields 30 (configuration, scenario) points; MR-002 is defined on 30 of the resulting relation comparisons and MR-004 on the remaining 24, giving 54 in all. The two agreed on 54 of them (1.00); the point set is a fixed enumeration, so no interval is attached. Both were written from the same reading of the semantics, so this bounds implementation accident only; it does not bound a shared blind spot, and we do not present it as corroboration.

## 8. Threats to validity

*Source versus execution* We read the predicates at source (built under JDK 8 and Maven; Section 7.1) and then executed the relations, unmodified apart from the build configuration and our harness (Section 7.2). The harness substitutes an HTTP observation channel—exact for the reachability predicate, not exercised for the content predicate, whose invocation count we report rather than claim cleanliness we did not demonstrate. The run converts one class of claim (the false proof is produced by the *published* relations on a real system); it does not convert the rest: the

scenario/configuration and cell studies remain statements about our instantiation, our exception vocabulary belongs to the same predicate family, and we report no sensitivity or specificity attributable to MST-wi.

*Metric and statistics* The specificity approximation uses two follow-up inputs per case—not the engine's count; the conclusion about the denominator's resolution does not depend on the exact value. Records within a design cell are not independent (Section 3) and the records are a fixed enumeration, so we report counts and rates rather than inferential intervals, and no result rests on statistical significance. The RQ2 ablation is paired and is reported as a discordance count.

*Two MRs, one family* We instantiated 2 MRs of one family—relations gated on the absence of evidence. The native run lifts the count to 10 of 11 catalog authorization relations but not the family: only the reachability-gated ones alarm in our run. This is the paper's single most important limitation; raising it requires relations gated otherwise, a different experiment.

*Real systems* Three systems (one lineage pair, one independent), chosen for accessibility, none vulnerable on the tested surface—so only the false-proof side is measured, within one domain, the git forge. Channel proportions are topology design parameters, not field prevalence. Constructed scenarios with construction-defined ground truth are why the real-system section exists; three systems tested once each are still not a population.

*Our own earlier claims* An earlier iteration proposed v3 as a mechanism on the strength of its precision (1.00), which was bought with 111 false negatives, and a control (`direct_only`, precision 0.50) that discarded the access result. Replacing it with a fair control produced the negative result. We record this as the condition any similar claim must meet: a control consuming the same observable quantity as the ground truth.

## 9. Related work

*Metamorphic security testing* Metamorphic testing was brought to web security by Mai et al. [3], who cast security test patterns—authentication, authorization, session—as metamorphic relations so that faults can be exposed without an oracle. The line matured into MST-wi [2], whose 76 system-agnostic relations over a Web-specific domain language are the object of our study. We propose no new relation and no better detector; we take the technique’s own admission that the role-to-resource mapping is “not always available” at face value and follow it to its consequence for what the adjudication can decide. A relation whose guard reads a crawl-time model is sound only while the model equals the authorization state, and that equality is exactly what the technique cannot check for itself.

*Detecting broken access control* A large and active literature detects access-control flaws directly, and it is worth separating the oracles it relies on. Black-box probing drives requests with and without the rights of another subject and looks for a missing refusal [4, 5], a line that runs back to AUTHSCOPE [6] and that has been extended to deployed systems under ethical constraints [7]. A complementary, white-box line recovers the *enforced* policy from the implementation—by reverse-engineering RBAC policies [8], by statically extracting authorization policies [9] or database access-control rules [10], or by locating the permission checks themselves [11]—and then flags violations against that recovered model: BOLA in database-backed applications [12], missing-owner-check and function-level flaws in Java web applications [13, 14], and privilege escalation in polyglot microservices [15]. Specification-driven API fuzzers generate the requests well but inherit the same oracle difficulty [16–18]; adding explicit access-policy violation checks [19] moves them toward the detection problem but does not remove the need for a ground truth, and a recent turn delegates generation or probing to language models [20–22]. Our question is orthogonal to all of these. Each reports detections against an oracle it holds privately—a policy, a seeded defect, a differential response; we ask what an existing oracle-substituting technique *decides* when the entitlement information it needs is not among its observables, and we make that a measurable property of the technique rather than of a system.

*Direct authorization oracles, and what is at stake here* Comparing responses across principals or privilege states is an established way to obtain an authorization oracle when the ground-truth relation is available, and the black-box and white-box lines above are instances of it. We propose no new direct oracle, and we do not claim that direct authorization comparison is novel. Our subject is a different question: what an existing oracle-substituting technique *decides* when it substitutes reachability for entitlement. The direct oracle

enters our design in exactly one role—as the fair, same-observable control that a necessity claim about the metamorphic step must face. Without such a control, a substitution can appear necessary only because the alternative it is compared against was not allowed to consume the same decision-relevant observables.

*LLMs and the oracle problem* LLMs are used for vulnerability detection and specification inference [23–26]; we use none in the detection loop. The transferable condition this paper exposes is that any claim of necessity must be tested against a fair substitute that consumes the same observable quantity as the ground truth.

*Measurement and evaluation* Our contribution is finally a measurement, and it connects to work on the quality of the artefacts security studies are built on—for instance the observation that vulnerability datasets are frequently mislabelled, duplicated or otherwise unfit for the claims drawn from them [27]. The artefact we report is of the same family but inside a study rather than a dataset: a specificity figure whose *denominator* is structurally unable to register a whole class of failure, and whose operating point silently assumes a reachability model equals the authorization state. We therefore read a testing technique’s own reported metric as a measurement instrument, stating explicitly what its denominator can and cannot resolve.

## 10. Conclusion

For object-level authorization, MST-wi’s predicates do not know whether the acting subject is entitled to the object; they know what the interface, as crawled, permitted. Instantiating them as defined produced three findings (Section 6).

*Finding 1 (temporal)* As designed, the adjudication false-proofs on legitimately granted access, deterministically (1.00 constructed; 0.4000 Gitea), because an out-of-band grant fits only the wrong one of the predicate’s two representations. On the real system the mismatch appears only under the crawl-time reading (10/48, vs. 0 live), with an endogenous control holding final visibility fixed: within this controlled design the only manipulated explanatory variable is when the model was read, not what was visible.

*Finding 2 (oracle)* The natural repair—a readable, sound authorization oracle that fully resolves the proposition under test—works, and in that regime the evaluated metamorphic step adds no adjudication information the oracle and the request do not already carry: a necessary condition (Lemma 1), whose empirical face is that a single ordinary request yields the same false-proof count (false proofs 0 vs. 0), and that of the 216 true-violation records the paired comparison disagrees on only 3, each time in favour of the plain request; no margin was pre-specified, so no inferential claim attaches. On Gitea the comparison is limited to false-proof behaviour, because no vulnerability was present on the tested surface.

**Finding 3 (measurement)** The reported specificity of 99.81% is not conditioned on an explicit check that the reachability model equals the authorization state, and its input-level denominator cannot provide stable scenario-level resolution for a scenario-uniform failure.

Support evidence: reading the implementation confirmed the semantics and added the mechanism—the crawl model is frozen by a cache never invalidated, and a missing crawl defaults to “unreachable”—and the positive control is a sanity check that closes the loop: on a seeded vulnerable target, for one fault pattern, the surface-gated adjudicator detects all planted instances with zero false positives.

**What we do not claim** We do not claim that MST-wi is wrong or say anything about its accuracy as a detector; the native run establishes a qualitative outcome, for one gate family, on one real system that on the tested endpoints and transitions enforces the expected decisions. Channel proportions are design parameters; our two MRs are two, of one family, and that is the ceiling on everything we say.

**What remains expensive** Moving from characterisation to something stronger needs, first, a real system that *contains* the flaw, so detection can be measured; and second, relations whose gate is not the absence of evidence—checkable against a subject’s entitlements without pretending that entitlements are reachable facts. The transferable condition—we call it the *same-observable condition*—is the one that killed our own mechanism: a component’s necessity must be tested against a substitute that consumes the *same* observable quantity as the ground truth; a substitute that consumes less does not establish incremental value. The MST-wi case study is an instance of that condition, not its scope: a component’s necessity is not established unless a fair control that lacks the component but consumes the same decision-relevant observables is empirically or formally shown to be weaker.

## Declarations

### Data availability

The raw records of the studies reported here are publicly available under the version-pinned snapshot artifact-v1 of the companion repository <https://github.com/ubfbuz3/JISA> (re-anchored to the submission commit; a byte-for-byte identical copy is archived at Zenodo, DOI: 10.5281/zenodo.XXXXXXX): (i) the scenario/configuration study (experiments/results/raw\_matrix.jsonl, ablation\_direct\*.json, alt\_matrix.jsonl); (ii) the real-system cell study across Gitea 1.22.6, Gogs 0.14.3 and GitLab 17.11.7 (experiments/real\_system/results/\*.json); (iii) the execution of the published relations on the native engine (experiments/native\_engine/results/\*.json); (iv) the seeded positive control (seeded.json). The analysis pipeline (analysis/\*.py: compute\_all.py and the ingest scripts) recomputes every number, table and figure in this paper from those records; running `bash reproduce.sh` regenerates all outputs and the PDF, and `bash verify_artifact.sh`

asserts the SHA-256 integrity of every raw record and byte-reproducible output. The manuscript L<sup>A</sup>T<sub>E</sub>X source ships in the same repository.

### Competing interests

The authors declare no competing interests, financial or otherwise.

### Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

### Authors’ contributions

Both authors contributed to the work. Qiulong Lin and Nanlei Chen jointly developed the conceptualisation and methodology, designed and implemented the software and analysis pipeline, carried out the investigation and formal analysis, and wrote the manuscript.

### Ethics

Not applicable. This study involved no human participants, no personal data, and no animals. All experiments were performed against software systems that we deployed and controlled locally.

## Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used Claude (Anthropic) to assist with literature verification, the development of the analysis scripts, and language editing. Where it contributed to analysis code, every generated or modified script was reviewed, executed, and its output checked by the authors against the raw records. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article. All experimental computations and final reported values were independently checked by the authors against the released scripts and raw records.

## References

- [1] OWASP Foundation, API1:2023 Broken Object Level Authorization, OWASP API Security Top 10, accessed 18 September 2026 (2023). URL <https://api-security.owasp.org/editions/2023/en/0xa1-broken-object-level-authorization/>
- [2] N. B. Chaleshtari, F. Pastore, A. Goknil, L. Briand, Metamorphic testing for web system security, IEEE Transactions on Software Engineering (2023) 3430–3471 doi:10.1109/tse.2023.3256322. URL <https://doi.org/10.1109/TSE.2023.3256322>
- [3] P. X. Mai, F. Pastore, A. Goknil, L. Briand, Metamorphic security testing for web systems, in: 2020 IEEE 13th International Conference on Software Testing, Validation and Verification (ICST), 2020, pp. 186–197. doi:10.1109/icst46399.2020.00028. URL <https://doi.org/10.1109/icst46399.2020.00028>
- [4] F. Liu, Y. Zhang, E. Li, W. Meng, Y.-Q. Shi, Q. Wang, M. Yang, Bacsan: Automatic black-box detection of broken-access-control vulnerabilities in web applications, in: Proceedings of the 2025 ACM SIGSAC Conference on Computer and Communications Security, 2025. doi:10.1145/3719027.3744825. URL <https://doi.org/10.1145/3719027.3744825>

- [5] I. P. A. Dharmadi, M. Alhanahnah, V.-T. Pham, F. Mohsen, F. Turkmen, Bacfuzz: Exposing the silence on broken access control vulnerabilities in web applications, *ArXiv* (2025). doi:10.48550/arxiv.2507.15984.  
URL <https://doi.org/10.48550/arXiv.2507.15984>
- [6] C. Zuo, Q. Zhao, Z. Lin, Authscope: Towards automatic discovery of vulnerable authorizations in online services, in: *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017. doi:10.1145/3133956.3134089.  
URL <https://doi.org/10.1145/3133956.3134089>
- [7] S. E. H. Chehade, F. Hantke, B. Stock, 403 forbidden? ethically evaluating broken access control in the wild, in: *2025 IEEE Symposium on Security and Privacy (SP)*, 2025, pp. 3218–3235. doi:10.1109/sp61157.2025.00252.  
URL <https://doi.org/10.1109/SP61157.2025.00252>
- [8] H. Le, L. K. Shar, D. Bianculli, L. Briand, D. C. Nguyen, Automated reverse engineering of role-based access control policies of web applications, *Journal of Systems and Software* 184 (2022) 111109. doi:10.1016/j.jss.2021.111109.  
URL <https://doi.org/10.1016/j.jss.2021.111109>
- [9] B. Berger, R. W. Nguemang, K. Sohr, R. Koschke, Static extraction of enforced authorization policies seeauthz, in: *2020 IEEE 20th International Working Conference on Source Code Analysis and Manipulation (SCAM)*, 2020, pp. 187–197. doi:10.1109/scam51674.2020.00026.  
URL <https://doi.org/10.1109/SCAM51674.2020.00026>
- [10] W. Zhang, D. Bali, J. Kerney, A. Panda, S. Shenker, Extracting database access-control policies from web applications, in: *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI '26)*, 2026, pp. 249–270.  
URL <https://www.usenix.org/conference/osdi26/presentation/zhang-wen>
- [11] Y.-Q. Shi, F. Liu, G. Yang, Y. Zhang, Y. Cao, E. Li, S. Chen, Facilitating access control vulnerability detection in modern java web applications with accurate permission check identification, *IEEE Transactions on Information Forensics and Security* (2025) 10683–10698. doi:10.1109/tifs.2025.3614424.  
URL <https://doi.org/10.1109/TIFS.2025.3614424>
- [12] Y. Huang, C. Shi, J. Lu, H. Li, H. Meng, L. Li, Detecting broken object-level authorization vulnerabilities in database-backed applications, in: *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*, 2024. doi:10.1145/3658644.3690227.  
URL <https://doi.org/10.1145/3658644.3690227>
- [13] F. Liu, Y.-Q. Shi, Y. Zhang, G. Yang, E. Li, M. Yang, Mocguard: Automatically detecting missing-owner-check vulnerabilities in java web applications, in: *2025 IEEE Symposium on Security and Privacy (SP)*, 2025, pp. 903–919. doi:10.1109/sp61157.2025.00010.  
URL <https://doi.org/10.1109/SP61157.2025.00010>
- [14] R. Feng, M. Zhang, A. Xia, J. Cheng, Y. Li, Y. Zhang, Y. Liu, Jauthguard: Automatic detection for broken access control in java web apis, in: *2025 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, 2025, pp. 3102–3107. doi:10.1109/smc58881.2025.11343625.  
URL <https://doi.org/10.1109/SMC58881.2025.11343625>
- [15] P. Li, H. Chong, Y. Cao, J. Yang, Detecting privilege escalation in polyglot microservices via agentic program analysis, in: *2026 IEEE Symposium on Security and Privacy (SP)*, 2026, pp. 1747–1765. doi:10.1109/sp63933.2026.00121.  
URL <https://doi.org/10.1109/SP63933.2026.00121>
- [16] V. Atlidakis, P. Godefroid, M. Polishchuk, Restler: Stateful rest api fuzzing, in: *2019 IEEE/ACM 41st International Conference on Software Engineering (ICSE)*, 2019, pp. 748–758. doi:10.1109/icse.2019.00083.  
URL <https://doi.org/10.1109/ICSE.2019.00083>
- [17] A. Arcuri, M. Zhang, J. P. Galeotti, Advanced white-box heuristics for search-based fuzzing of rest apis, *ACM Transactions on Software Engineering and Methodology* (2024). doi:10.1145/3652157.  
URL <https://doi.org/10.1145/3652157>
- [18] S. Seran, M. Zhang, O. Duman, A. Arcuri, Handling web service interactions in fuzzing with search-based mock-generation, *ACM Transactions on Software Engineering and Methodology* (2026). doi:10.1145/3731558.  
URL <https://doi.org/10.1145/3731558>
- [19] Ömür Şahin, M. Zhang, A. Arcuri, Enhancing rest api fuzzing with access policy violation checks and injection attacks, *Journal of Systems and Software* 243 (2027). doi:10.1016/j.jss.2026.113060.  
URL <https://doi.org/10.1016/j.jss.2026.113060>
- [20] M. A. Uddin, S. Weerasinghe, C. Wojtak, T. Cerný, D. Silva-Junior, M. E. S. Ribeiro, A. S. Abdelfattah, Automated generation of microservice authorization tests using large language models, in: *2026 International Conference on Service-Oriented System Engineering (SOSE)*, 2026, pp. 41–50. doi:10.1109/sose71128.2026.00014.  
URL <https://doi.org/10.1109/SOSE71128.2026.00014>
- [21] J. Sun, R. Tian, S. Zhao, X. Zhou, H. Wang, Y. Zhang, Llm-assisted idor detection in hospital mini-programs: Risks to pii and phi, in: *2025 IEEE 24th International Conference on Trust, Security and Privacy in Computing and Communications (TrustCom)*, 2025, pp. 2078–2086. doi:10.1109/trustcom66490.2025.00242.  
URL <https://doi.org/10.1109/Trustcom66490.2025.00242>
- [22] E. Wang, W. Xie, S. Li, R. Liu, Y. Zhou, Z. Wang, B. Wang, Large language model-powered protected interface evasion: Automated discovery of broken access control vulnerabilities in internet of things devices, *Sensors (Basel, Switzerland)* (2025). doi:10.3390/s25092913.  
URL <https://doi.org/10.3390/s25092913>
- [23] H. Wang, Z. Wang, P. Liu, A hybrid llm workflow can help identify user privilege related variables in programs of any size, *ArXiv* (2024). doi:10.48550/arxiv.2403.15723.  
URL <https://doi.org/10.48550/arXiv.2403.15723>
- [24] C. Wen, J. Cao, J. Su, Z. Xu, S. Qin, M. He, C. Tian, Enchanting program specification synthesis by large language models using static analysis and program verification, in: *International Conference on Computer Aided Verification*, 2024, pp. 302–328. doi:10.1007/978-3-031-65630-9\_16.  
URL [https://doi.org/10.1007/978-3-031-65630-9\\_16](https://doi.org/10.1007/978-3-031-65630-9_16)
- [25] R. Wang, H. Zhuang, H. Cao, S. Chen, Measuring the reasoning boundaries of large language models for implicit security invariants in code: A controlled empirical study, in: *2026 IEEE 50th Annual Computers, Software, and Applications Conference (COMPSAC)*, 2026, pp. 3146–3157. doi:10.1109/compsac69091.2026.00467.  
URL <https://doi.org/10.1109/COMPSAC69091.2026.00467>
- [26] T. Le-Cong, B. Le, T. Murray, Can llms reason about program semantics? a comprehensive evaluation of llms on formal specification inference, in: *Annual Meeting of the Association for Computational Linguistics*, 2025, pp. 21991–22014. doi:10.18653/v1/2025.acl-long.1068.  
URL <https://aclanthology.org/2025.acl-long.1068/>
- [27] R. Croft, M. A. Babar, M. M. Kholoosi, Data quality for software vulnerability datasets, in: *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*, 2023, pp. 121–133. doi:10.1109/icse48619.2023.00022.  
URL <https://doi.org/10.1109/ICSE48619.2023.00022>